



인공지능의 기초

마르코프 결정 과정

Markov Decision Processes

⚙️ A Markov decision process is an MRP with decisions:

$\langle S, \mathbf{A}, P, R, \gamma \rangle$

- ▶ A set of states $S = \{s_1, s_2, \dots, s_n\}$
- ▶ A set of actions $\mathbf{A} = \{a_1, a_2, \dots, a_m\}$
- ▶ Transition function $P: S \times \mathbf{A} \rightarrow S$, $P_{ss'}^a = \mathbb{P}[S_{t+1} = s' | S_t = s, A_t = a]$
- ▶ Reward function $R: S \times \mathbf{A} \rightarrow \mathbb{R}$, $R_s^a = \mathbb{E}[R_{t+1} | S_t = s, A_t = a]$
- ▶ Discount factor $\gamma \in [0, 1]$

MEMO

Markov Decision Processes

⚙️ A policy π is a distribution over actions given states

$$\pi(a|s) = P[A_t = a | S_t = s]$$

- ▶ MDP policies depend on the current state (not the history)
- ▶ Policies are stationary (time-independent) $A_t \sim \pi(\cdot | S_t), \forall t > 0$

MEMO

Value Functions

- ⊗ The state-value function $V_{\pi}(s)$ is the expected return starting from state s , **under a policy π**

$$V_{\pi}(s) = \mathbb{E}_{\pi} [G_t | S_t = s]$$

- ⊗ The action-value function $Q_{\pi}(s, a)$ is the expected return starting from state s , taking action a , under a policy π

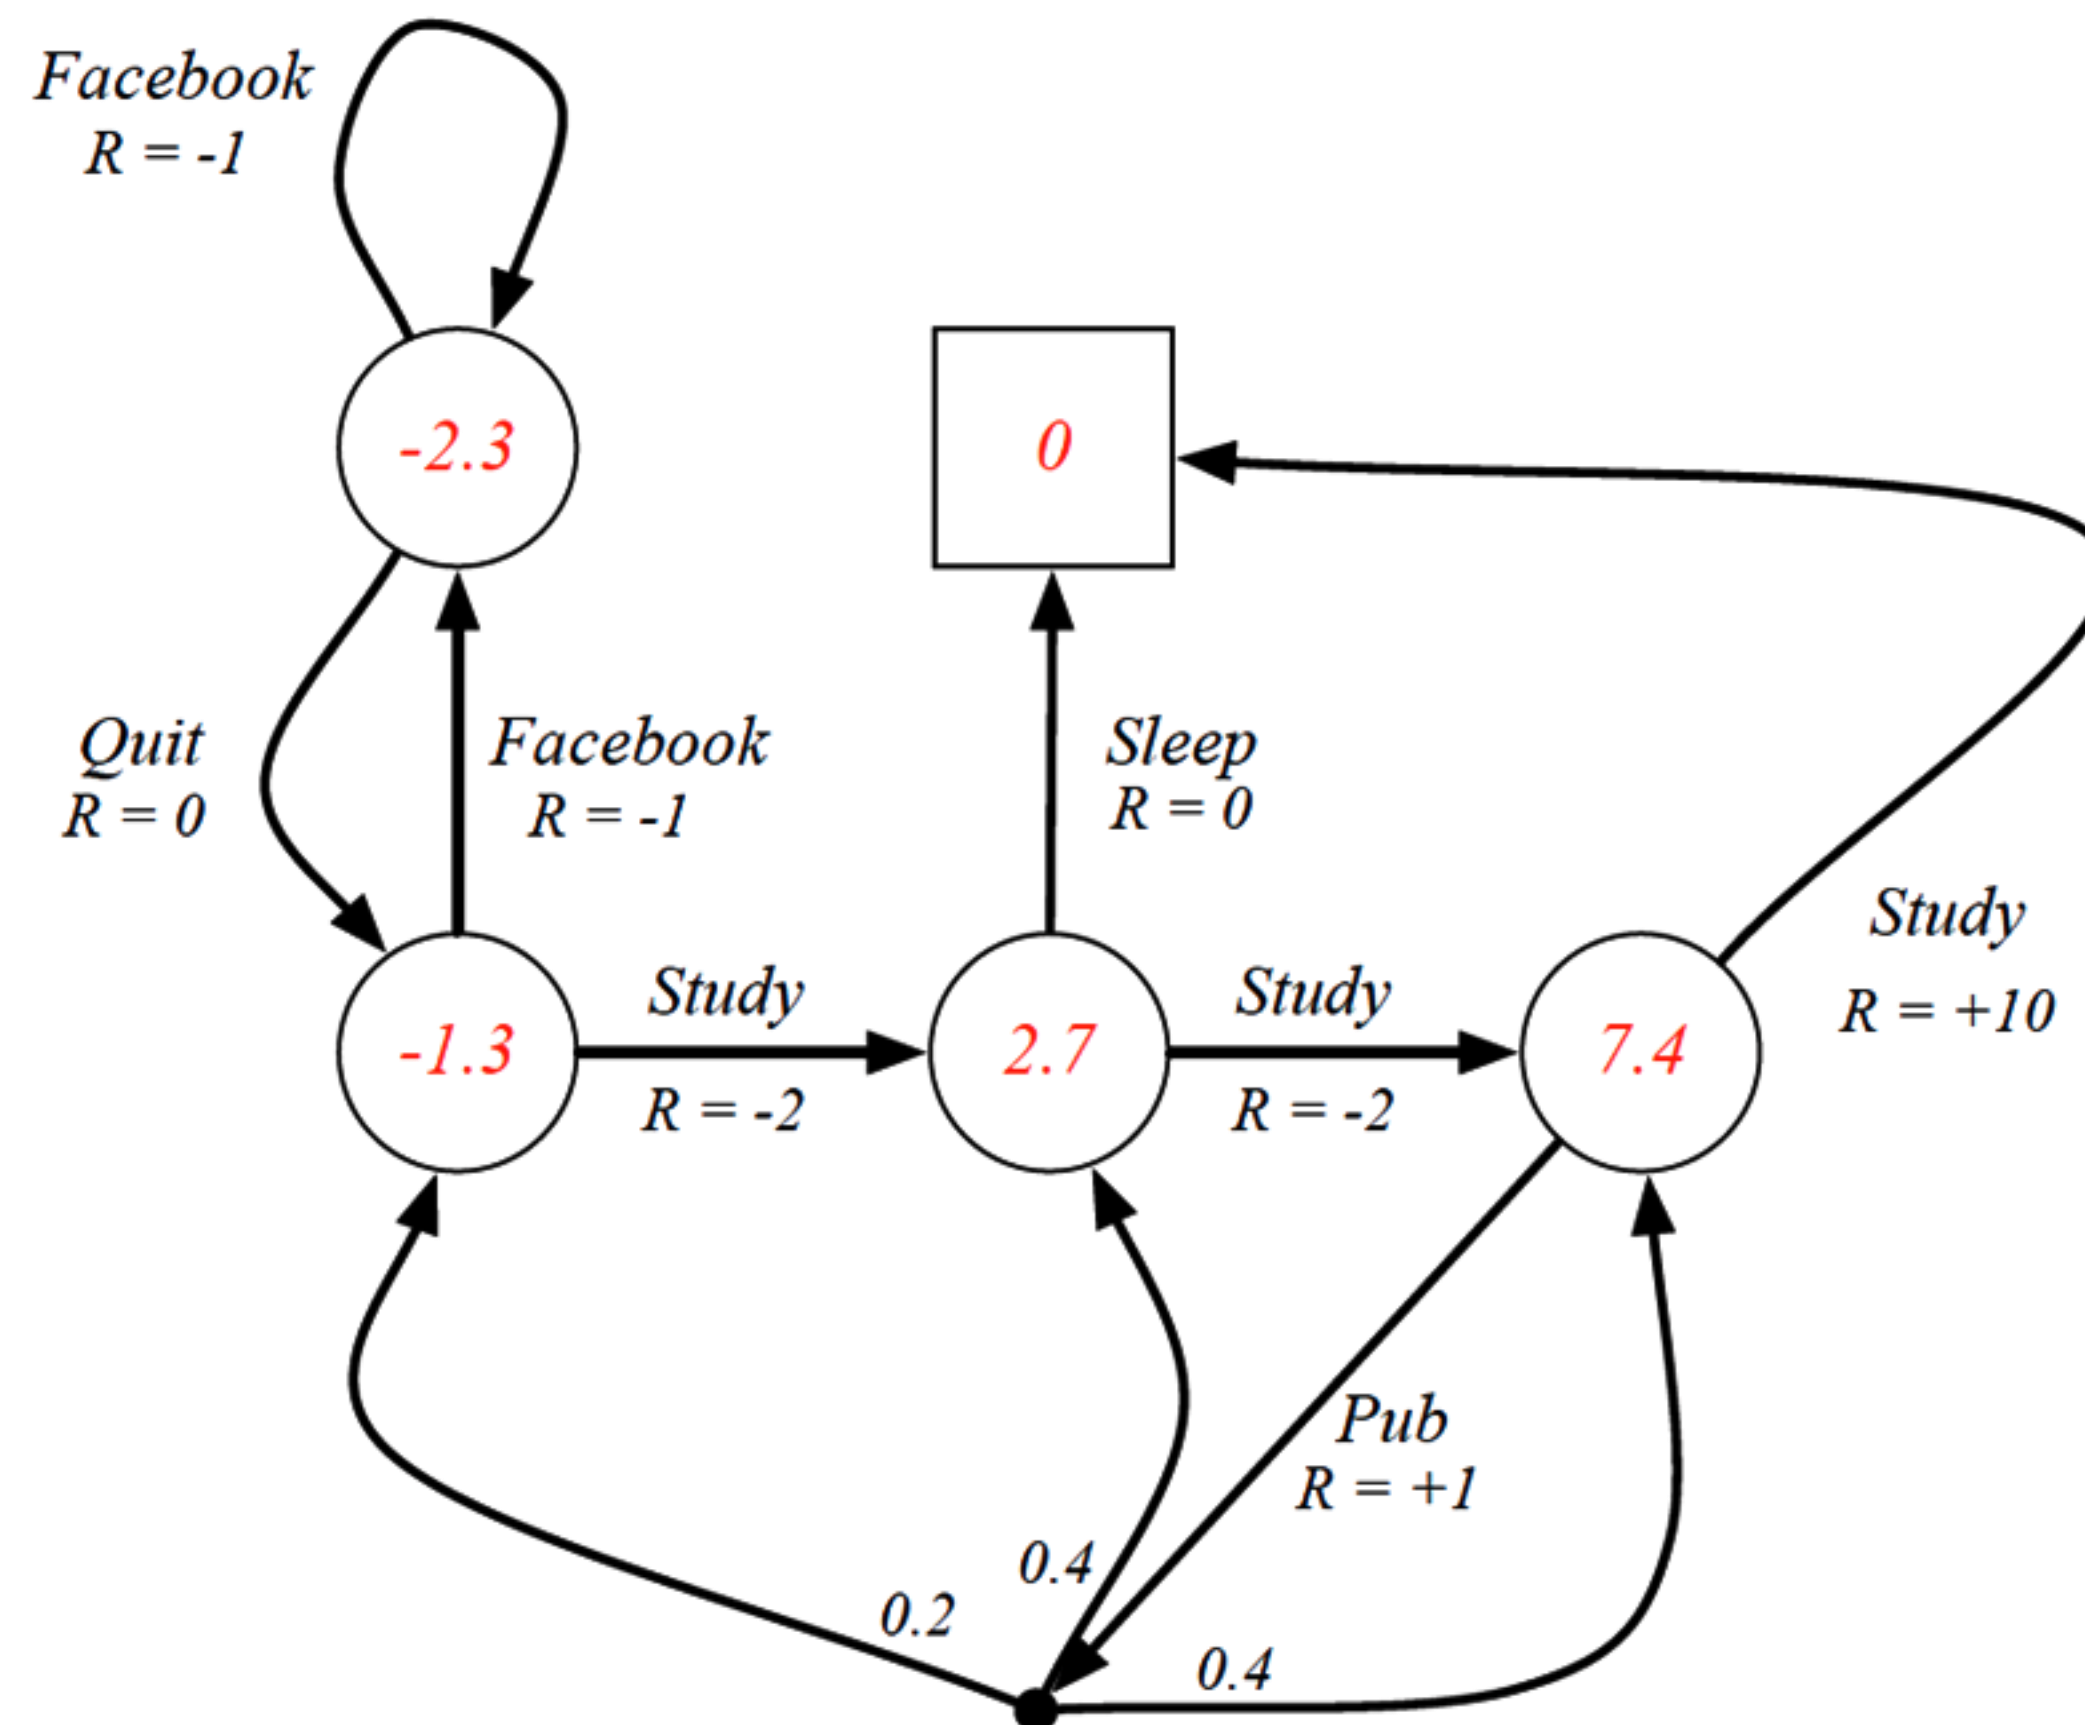
$$Q_{\pi}(s, a) = \mathbb{E}_{\pi} [G_t | S_t = s, A_t = a]$$

MEMO

Example of Student MDP

❗ Random policy with $\gamma = 1$

► $V_{\pi}(s)$ for $\pi(a|s) = 0.5$



MEMO

Bellman **Expectation** Equation for MDPs

⚙ The value function can be decomposed into two parts:

- ▶ Immediate reward R_{t+1}
- ▶ Discounted value of successor state $\gamma V(s_{t+1})$

⚙ The state-value function can be decomposed

$$\begin{aligned} V_{\pi}(s) &= \mathbb{E}_{\pi}[G_t | s_t = s] \\ &= \mathbb{E}_{\pi}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots | s_t = s] \\ &= \mathbb{E}_{\pi}[R_{t+1} + \gamma(R_{t+2} + \gamma R_{t+3} + \cdots) | s_t = s] \\ &= \mathbb{E}_{\pi}[R_{t+1} + \gamma G_{t+1} | s_t = s] \\ &= \mathbb{E}_{\pi}[R_{t+1} + \gamma V_{\pi}(s_{t+1}) | s_t = s] \end{aligned}$$

MEMO

Bellman **Expectation** Equation for MDPs

⚙ The value function can be decomposed into two parts:

- ▶ Immediate reward R_{t+1}
- ▶ Discounted value of successor state $\gamma V(s_{t+1})$

⚙ The state-value function can be decomposed

$$V_{\pi}(s) = \mathbb{E}_{\pi} [R_{t+1} + \gamma V(S_{t+1}) | S_t = s]$$

⚙ The action-value function can be decomposed

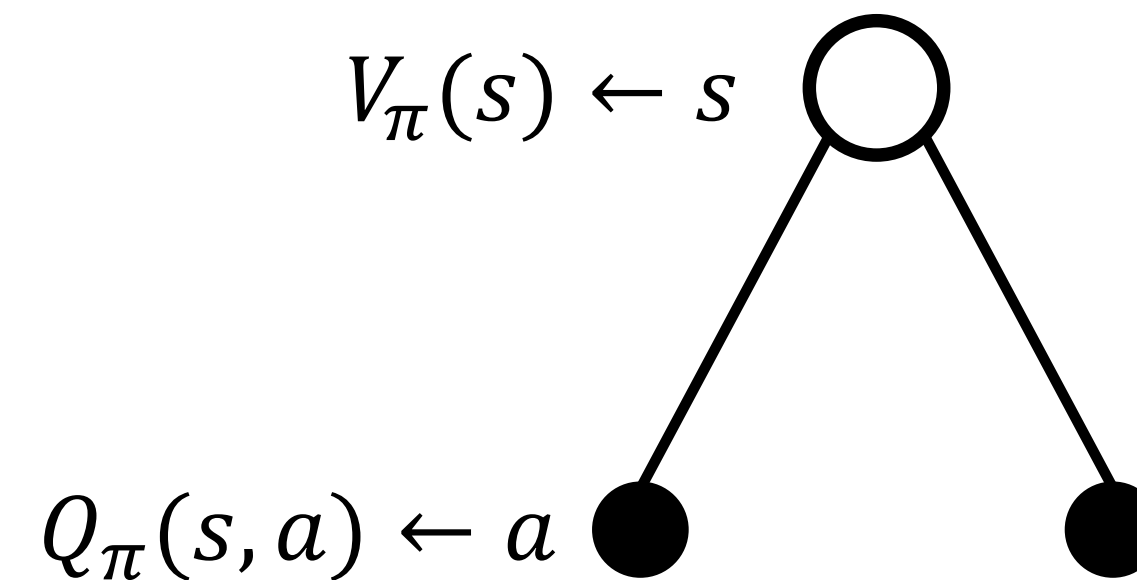
$$Q_{\pi}(s, a) = \mathbb{E}_{\pi} [R_{t+1} + \gamma Q_{\pi}(S_{t+1}, A_{t+1}) | S_t = s, A_t = a]$$

MEMO

Bellman Equation for V_π and Q_π

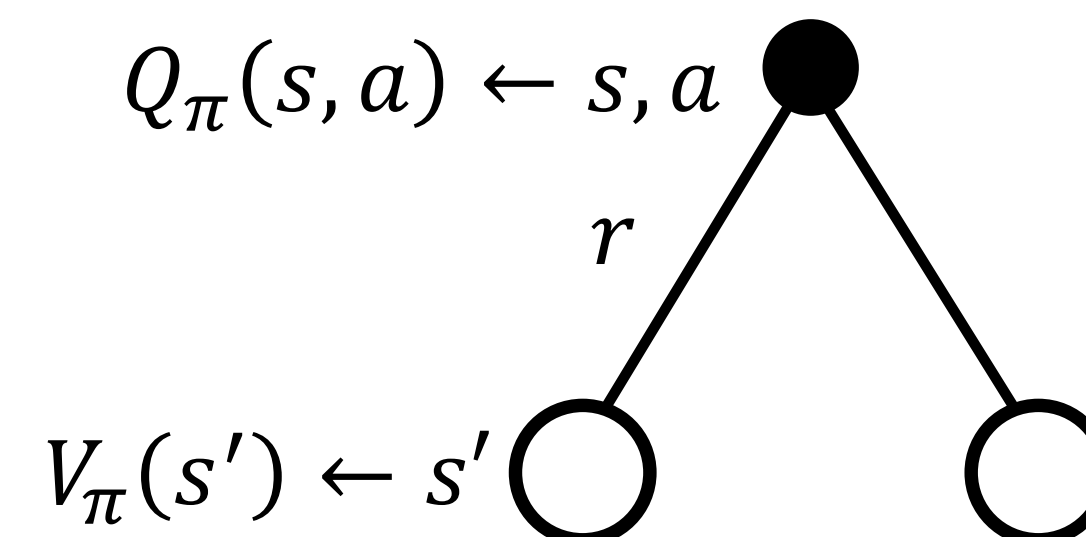
⚙ Bellman expectation equation for V_π

$$V_\pi(s) = \sum_{a \in A} \pi(a|s) Q_\pi(s, a)$$



⚙ Bellman expectation equation for Q_π

$$Q_\pi(s, a) = R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a V_\pi(s')$$

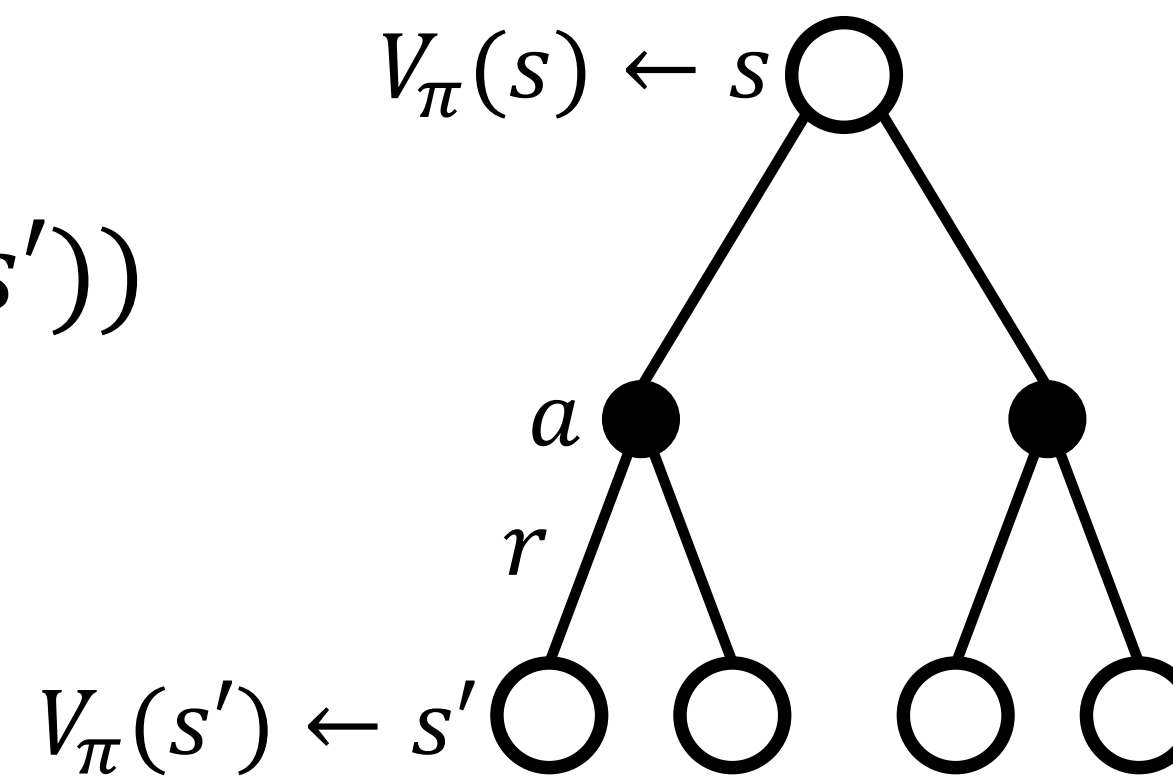


MEMO

Bellman Equation for V_π and Q_π

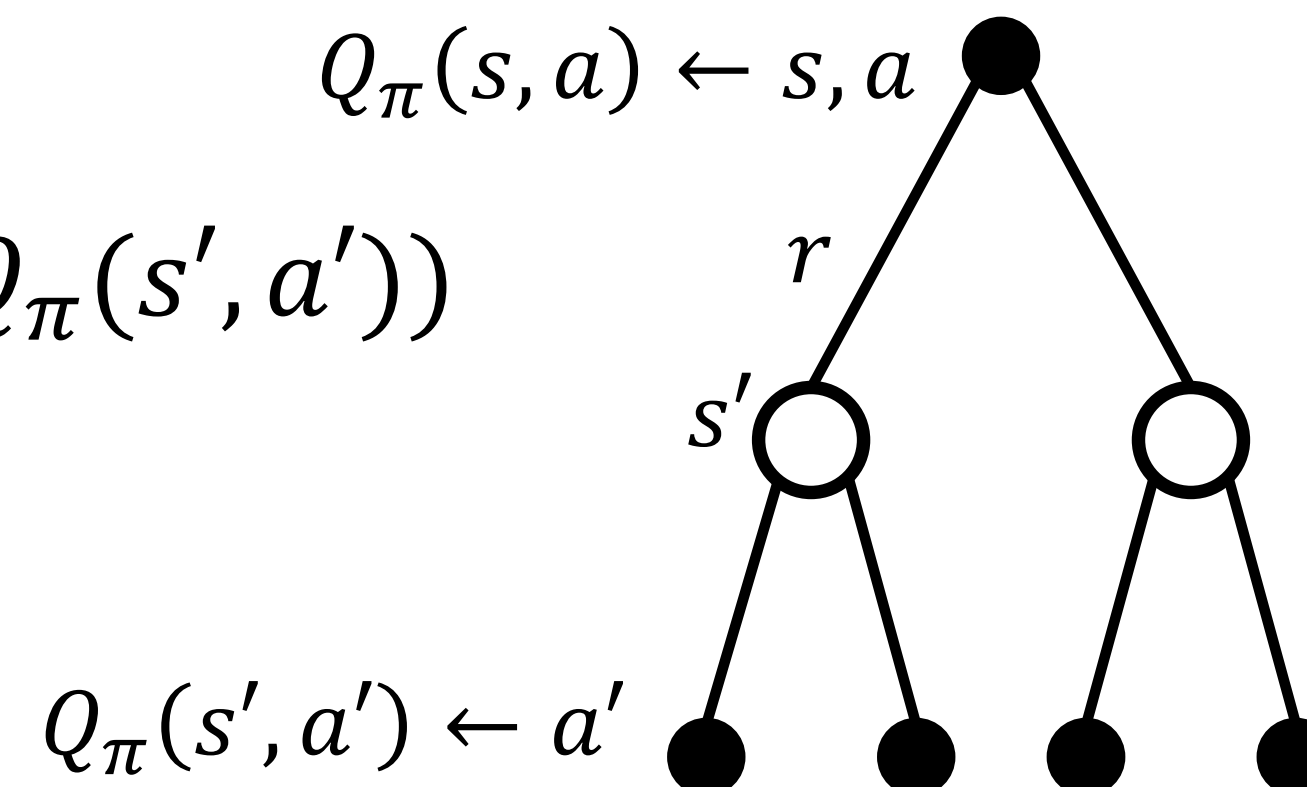
⚙ Bellman expectation equation for V_π (2)

$$V_\pi(s) = \sum_{a \in A} \pi(a|s) (R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a V_\pi(s'))$$



⚙ Bellman expectation equation for Q_π (2)

$$Q_\pi(s, a) = R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a \left(\sum_{a' \in A} \pi(a'|s') Q_\pi(s', a') \right)$$

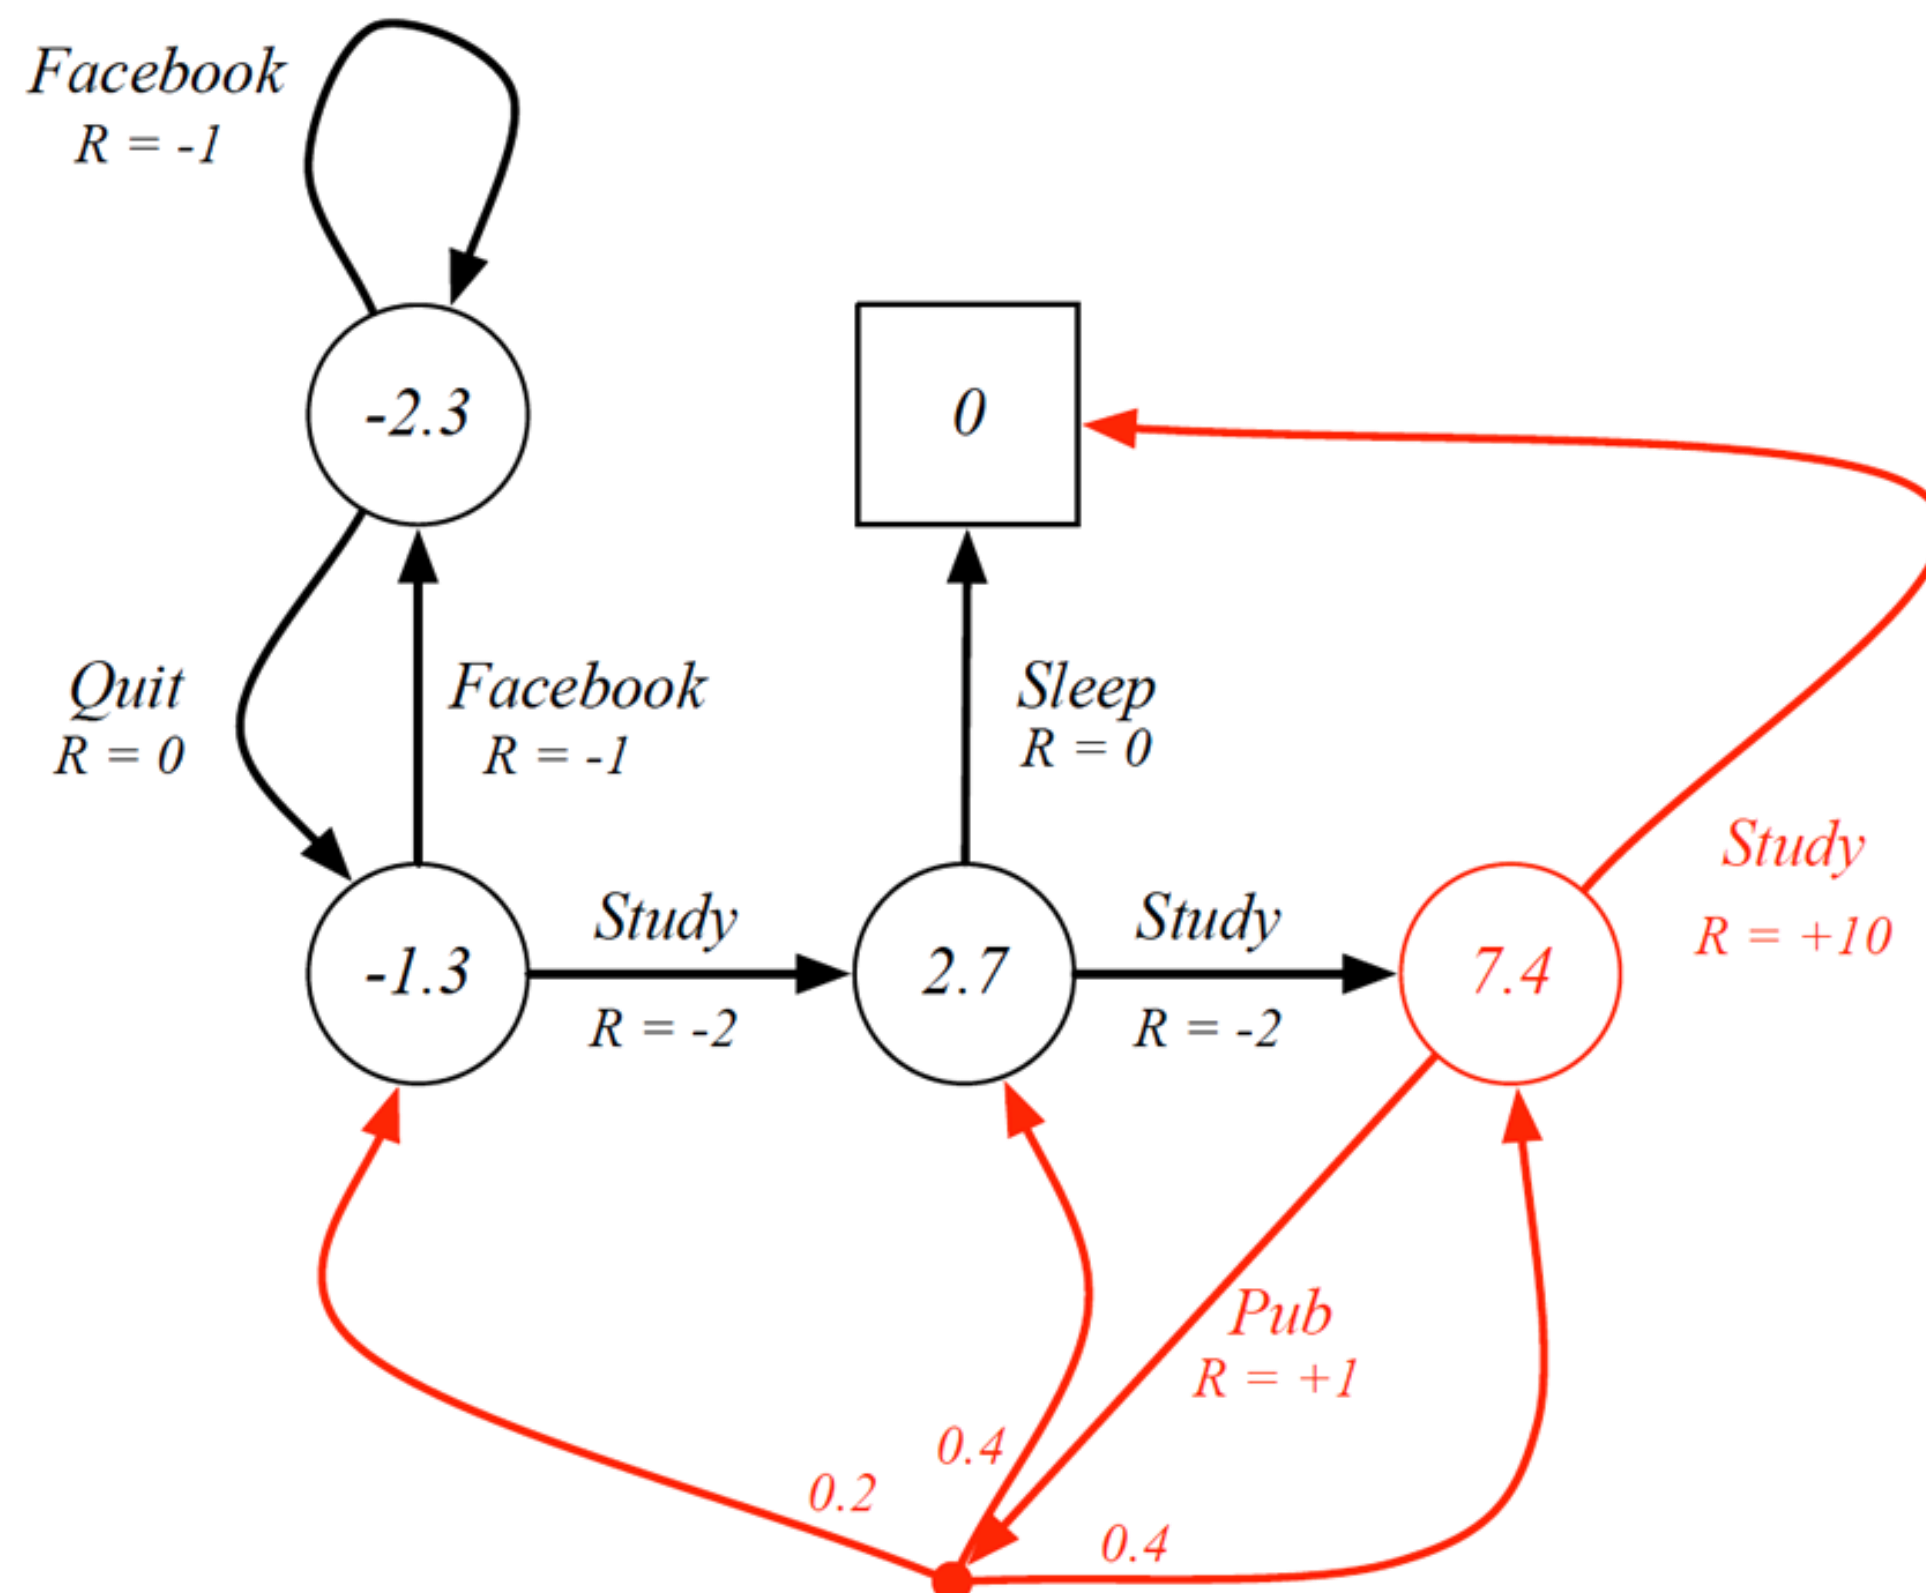


MEMO

Bellman Equation in Student MDP

$$V_{\pi}(s) = \sum_{a \in A} \pi(a|s) (R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a V_{\pi}(s'))$$

$$V_{\pi}(s) = 0.5 * (1 + 0.2 * (-1.3) + 0.4 * 2.7 + 0.4 * 7.4) + 0.5 * 10 = 7.4$$



MEMO

Bellman Expectation Equation (Matrix Form)

- ⊗ Can be expressed concisely in a matrix form

$$V_{\pi} = R^{\pi} + \gamma P^{\pi} V_{\pi}$$

$$\begin{bmatrix} V_{\pi}(s_1) \\ \dots \\ V_{\pi}(s_n) \end{bmatrix} = \begin{bmatrix} R_1^{\pi} \\ \dots \\ R_n^{\pi} \end{bmatrix} + \gamma \begin{bmatrix} P_{11}^{\pi} & \dots & P_{1n}^{\pi} \\ \dots & \dots & \dots \\ P_{n1}^{\pi} & \dots & P_{nn}^{\pi} \end{bmatrix} \begin{bmatrix} V_{\pi}(s_1) \\ \dots \\ V_{\pi}(s_n) \end{bmatrix}$$

- ⊗ It is a linear equation, so solved by

$$V_{\pi} = (I - \gamma P^{\pi})^{-1} R^{\pi}$$

- ▶ Computational complexity is $O(n^3)$ for n states
- ▶ Other approach? Dynamic programming, Monte-Carlo evaluation, Temporal-Difference learning

MEMO

Optimal Value Functions

- ⚙ The optimal state value function $V_*(s)$ is the max value function over all policies

$$V_*(s) = \max_{\pi} V_{\pi}(s)$$

- ⚙ The optimal action-value function $Q_*(s, a)$ is the maximum action-value function over all policies

$$Q_*(s, a) = \max_{\pi} Q_{\pi}(s, a)$$

- ⚙ Theorems: For any MDP

- ▶ There exists an optimal policy $\pi_* \geq \pi, \forall \pi$
- ▶ All optimal policies achieve the optimal state-value, $V_{\pi_*}(s) = V_*(s)$
- ▶ All optimal policies achieve the optimal action-value,
 $Q_{\pi_*}(s, a) = Q_*(s, a)$

MEMO

Finding an Optimal Policy

⚙️ An optimal policy can be found by maximizing over $Q_*(s, a)$

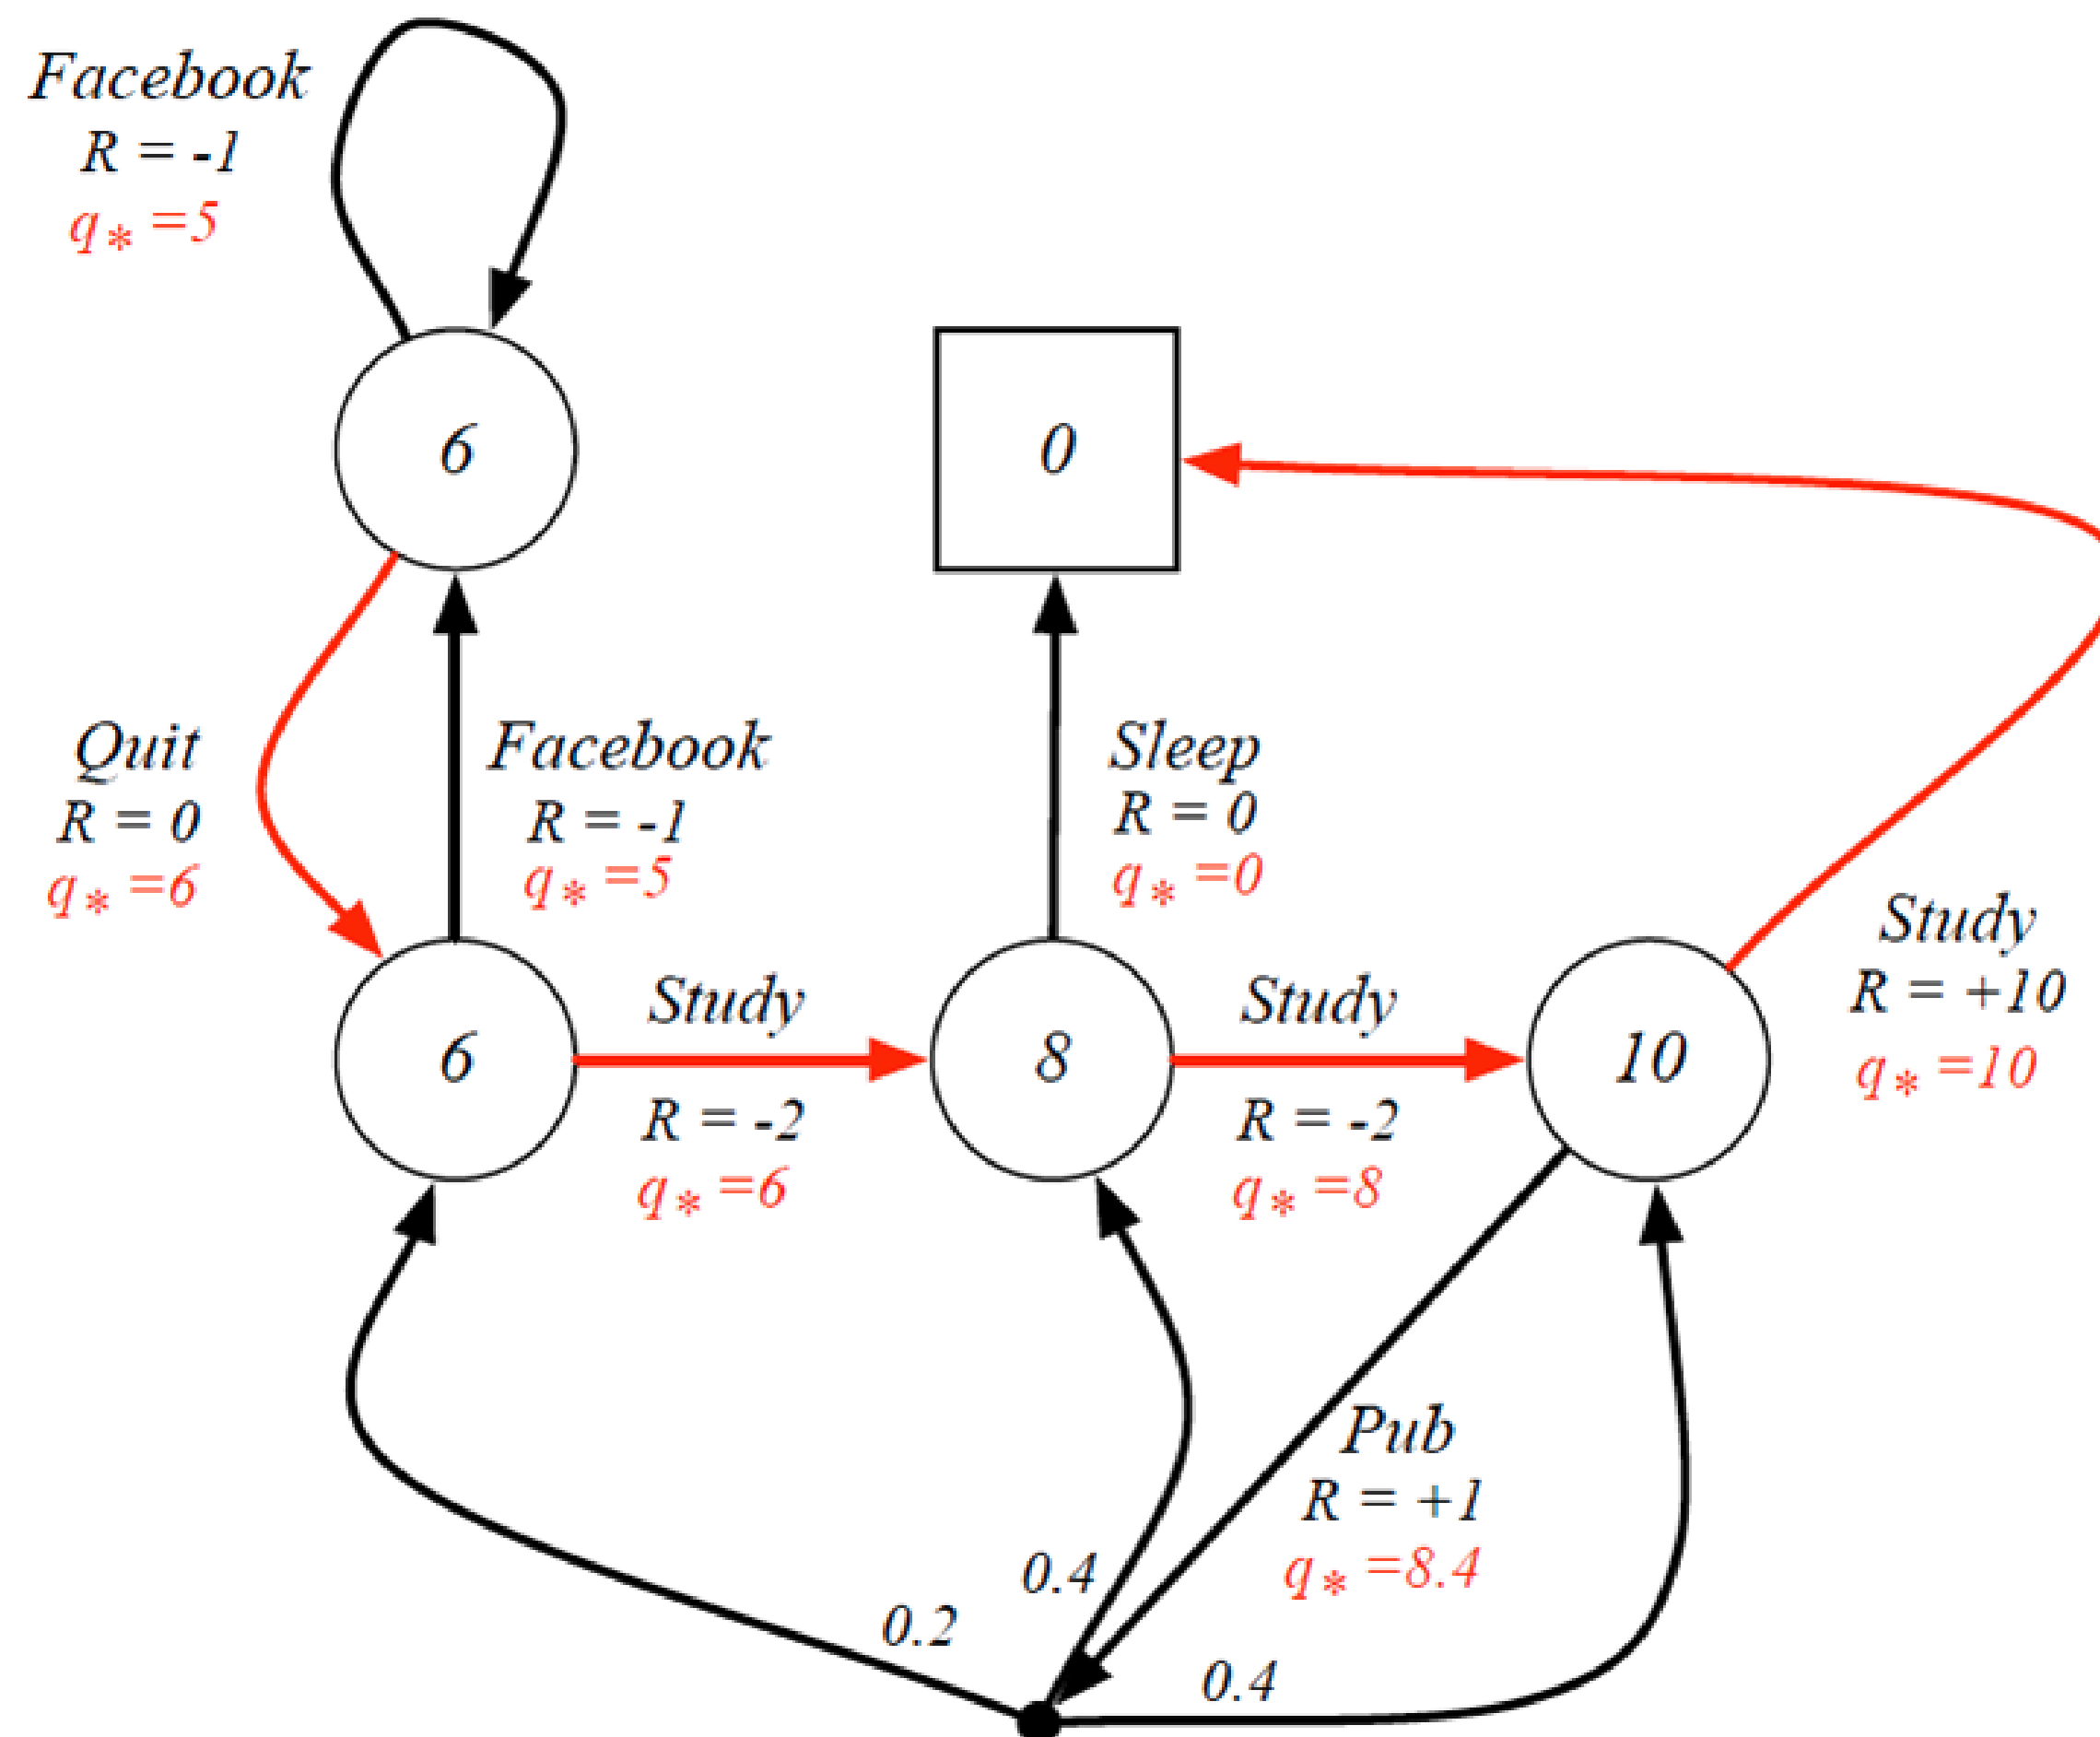
$$\pi_*(a|s) = \begin{cases} 1 & \text{if } a = \operatorname{argmax}_{a \in A} Q_*(s, a) \\ 0 & \text{otherwise} \end{cases}$$

► If we know $Q_*(s, a)$, we immediately have the optimal policy

MEMO

Optimal Policy for Student MDP

$$\pi_*(a|s) \text{ for } \gamma = 1$$



MEMO



인공지능의 기초

동적 프로그래밍을 통한 마르코프 결정 과정

Dynamic Programming

⚙️ A very general solution method for problems which have two properties

1. Optimal substructure

- ▶ Optimal solution can be decomposed into subproblems

2. Overlapping subproblems

- ▶ Subproblems recur many times
- ▶ Solutions can be cached and reused

⚙️ Markov decision processes satisfy both properties

- ▶ Bellman equation gives recursive decomposition
- ▶ Value function stores and reuses solutions

MEMO

Prediction and Control

⚙ Prediction: evaluate the future

- ▶ Given an MDP $\langle S, A, P, R, \gamma \rangle$ and a policy π
- ▶ Output: a value function V_π

Iterative policy evaluation!

⚙ Control: optimize the future

- ▶ Given an MDP $\langle S, A, P, R, \gamma \rangle$
- ▶ Output: optimal policy π_* (and optimal value function V_*)

MEMO

Iterative Policy Evaluation

- ❗ Problem: evaluate a given policy π
- ❗ Solution: iteratively apply Bellman expectation backup
 - ▶ Converge to a real V_π ($V_1 \rightarrow V_2 \rightarrow \dots \rightarrow V_\pi$)
 - ▶ At each iteration $k + 1$, for all states $s \in S$, update $V_{k+1}(s)$ from $V_k(s')$ where s' is a successor state of s
- ❗ Iteratively compute until convergence

$$\mathbf{V}^{k+1} = \mathbf{R}^\pi + \gamma \mathbf{P}^\pi \mathbf{V}^k$$

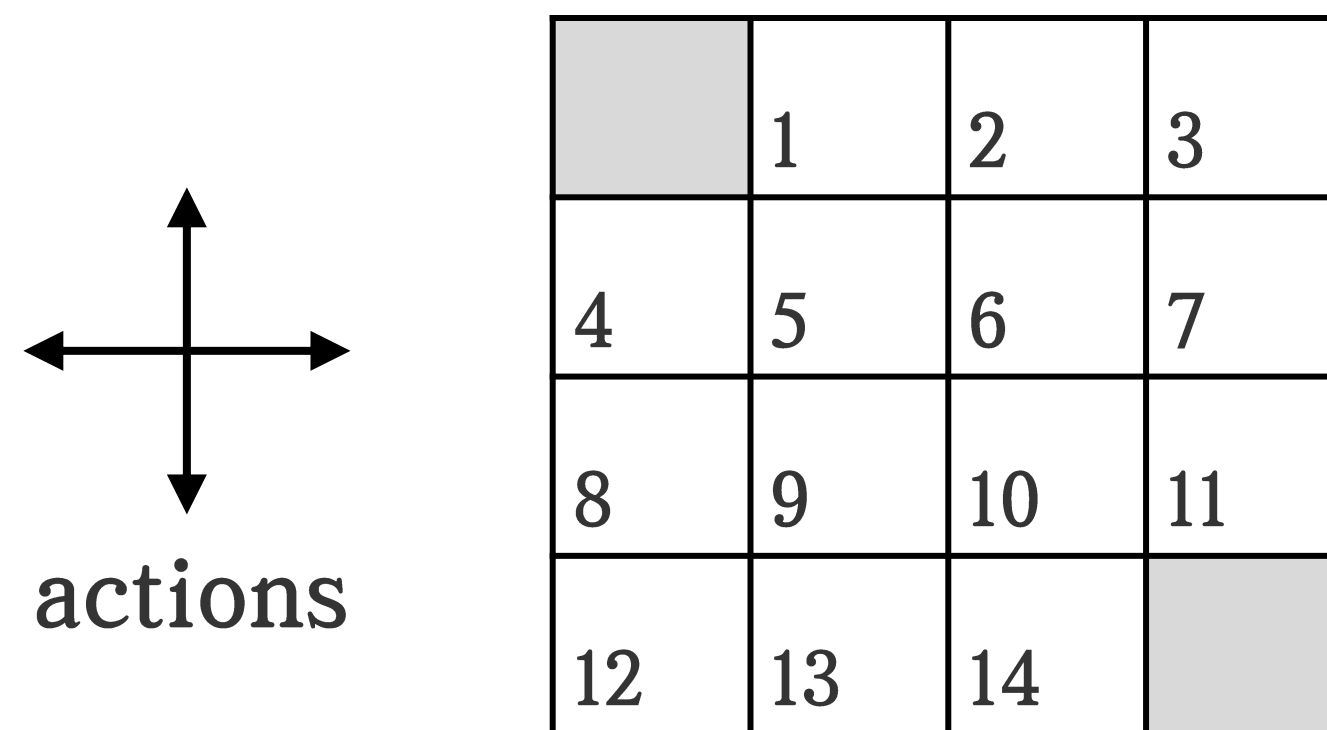
- ▶ Matrix form of Bellman expectation equation

$$V_\pi(s) = \sum_{a \in A} \pi(a|s) (R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a V_\pi(s'))$$

MEMO

Evaluating Random Policy in Small Gridworld

⚙️ Problem setup



- ▶ Undiscounted episodic MDP ($\gamma = 1$)
- ▶ Terminal state: two shaded squares
- ▶ Actions leading out of the grid leave state unchanged
- ▶ Reward is -1 until the terminal state is reached
- ▶ Agent follows uniform random policy

$$\pi(n | \cdot) = \pi(e | \cdot) = \pi(s | \cdot) = \pi(w | \cdot) = 0.25$$

MEMO

Iterative Policy Estimation in Small Gridworld

Problem setup

V^k for the
random policy

Greedy policy
w.r.t. V^k

Converged
optimal policy

$k = 0$

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

	↕	↕	↕
↕	↕	↕	↕
↕	↕	↕	↕
↕	↕	↕	

$k = 3$

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0

	←	←	↘
↑	↖	↘	↓
↑	↗	↘	↓
↖	→	→	

$k = 1$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

	←	↕	↕
↑	↕	↕	↕
↕	↕	↕	↓
↕	↕	→	

$k = 10$

0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0

	←	←	↘
↑	↖	↘	↓
↑	↗	↘	↓
↖	→	→	

$k = 2$

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

	←	←	↕
↑	↖	↕	↓
↑	↕	↘	↓
↕	→	→	

$k = \dots$

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0

	←	←	↘
↑	↖	↘	↓
↑	↗	↘	↓
↖	→	→	

MEMO

Policy Improvement

- ⚙️ Now we know how to evaluate V_π for a given policy π
- ⚙️ How to improve a policy?
 - ▶ Initialize a policy π
 - ▶ Evaluate the policy π to compute V_π
 - ▶ Improve the policy by acting greedily with respect to V_π

$$\pi' = \text{greedy}(V_\pi)$$

$$\pi'(s) = \operatorname{argmax}_{a \in A} Q_\pi(s, a)$$

- ▶ This process of policy iteration always converges to π_*

MEMO

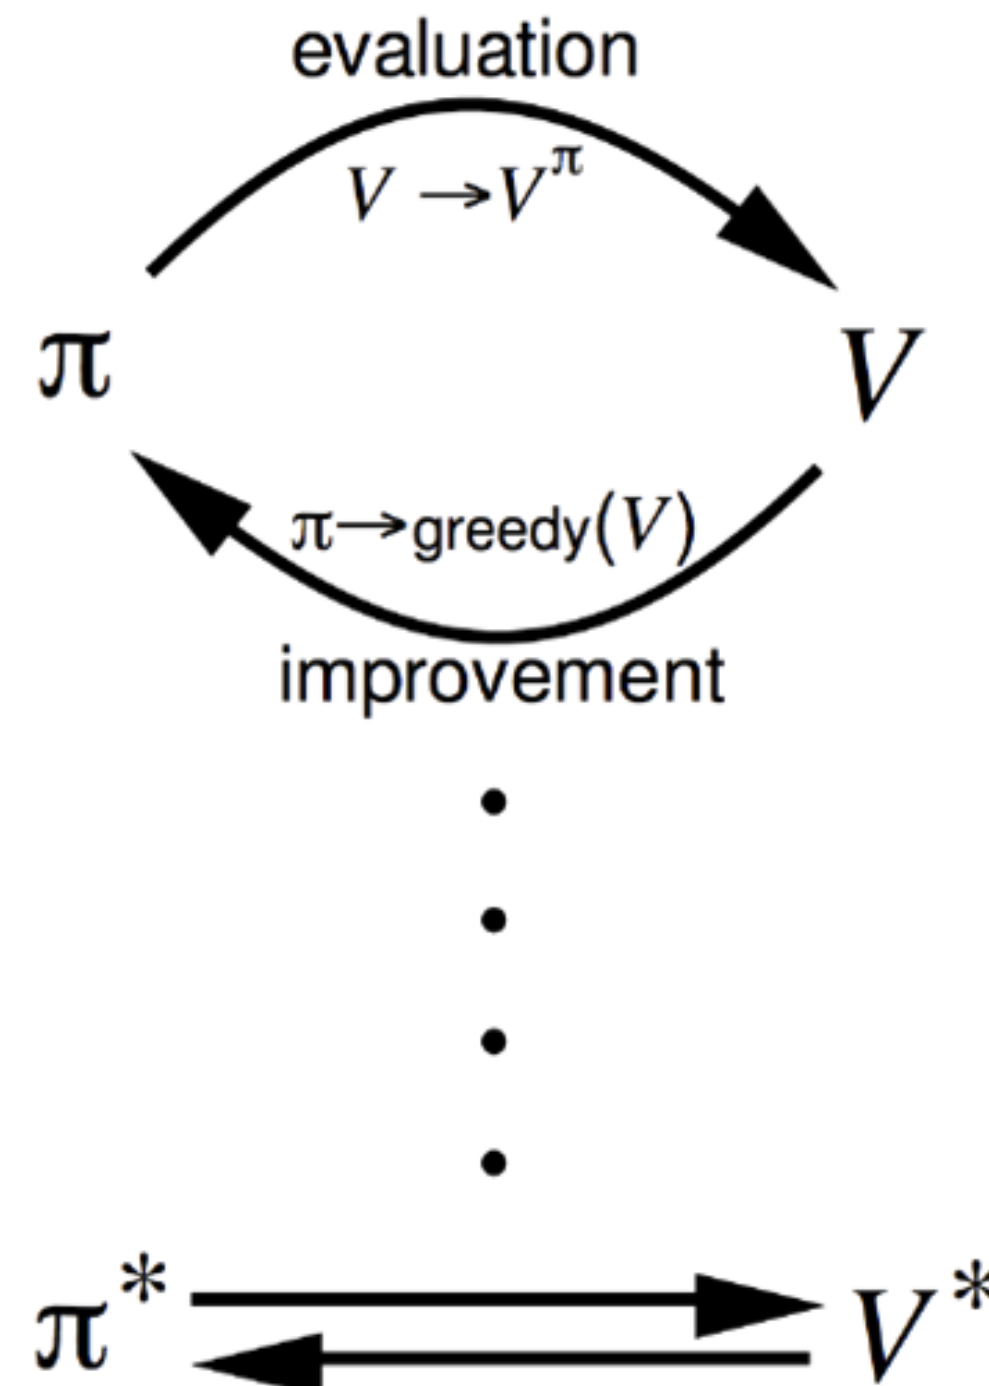
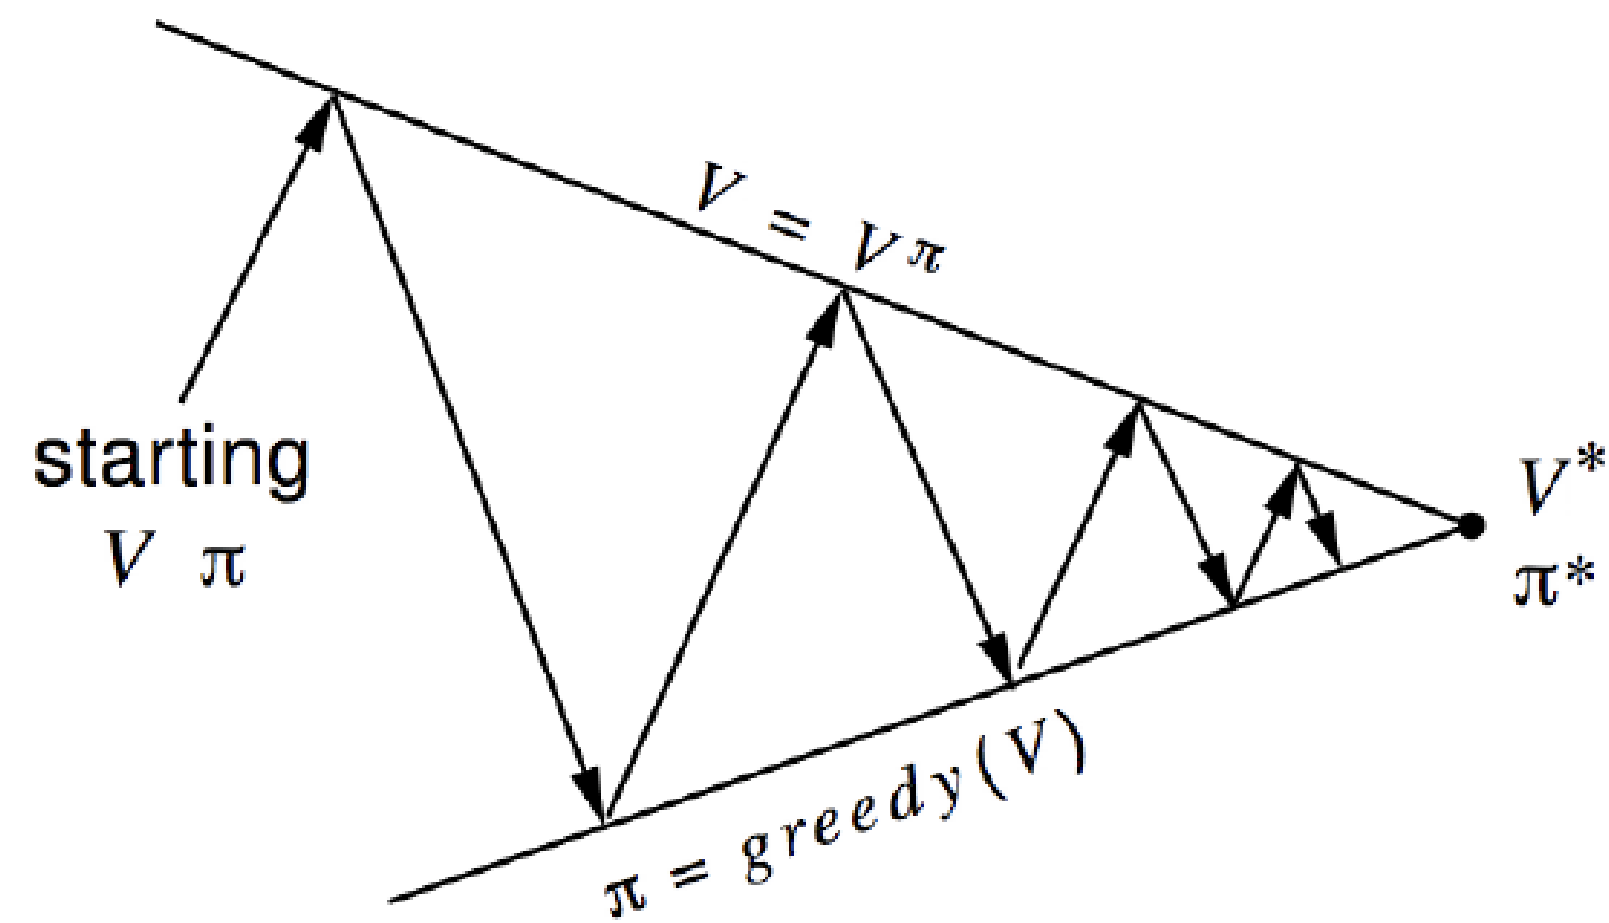
Policy Iteration

⚙️ Policy evaluation: estimate V_π

► Iterative policy evaluation

⚙️ Policy improvement: generate $\pi' \geq \pi$

► Greedy policy improvement



MEMO

DP Algorithms

Problem	Bellman equation	Algorithm
Prediction	Bellman expectation equation	Iterative policy evaluation
Control	Bellman expectation equation + Greedy policy improvement	Policy iteration

- ▶ Algorithms are based on state-value function $V_{\pi}(s)$ or $V_*(s)$
- ▶ Complexity $O(mn^2)$ per iteration, for m actions and n states
- ▶ Could also apply to action-value function $Q_{\pi}(s, a)$ or $Q_*(s, a)$
- ▶ Complexity $O(m^2n^2)$ per iteration

MEMO